

Metryki RED dla aplikacji REST z Prometheus + Grafana + AlertManager



WOJCIECH BARCZYŃSKI (WOJCIECH.BARCZYNSKI@SMACC.IO)

WOJCIECH BARCZYŃSKI


- Lead Software Developer - SMACC (FinTech/AI)
- Before:
System Engineer i Developer Lyke
- Before:
1000+ nodes, 20 data centers with Openstack
- Interests:
Working software, Effective and Satisfied Teams

WHY?
MONOLIT ;)



WHY?

MICROSERVICES ;)



	Monitoring	Logging	Tracing
Setup	Easy	Diff	Diff
TCO	Low	Very High	High
Debugging	Low	High	High
Detecting	High	Low	Low

NOT A SILVER-BULLET

but:

- Easy to setup
- Immediately value

Suprisengly: the last one implemented

CENTRALIZED LOGGING

- Usually much too late
- Post-mortem
- Hard to find the needle
- Like a debugging vs testing

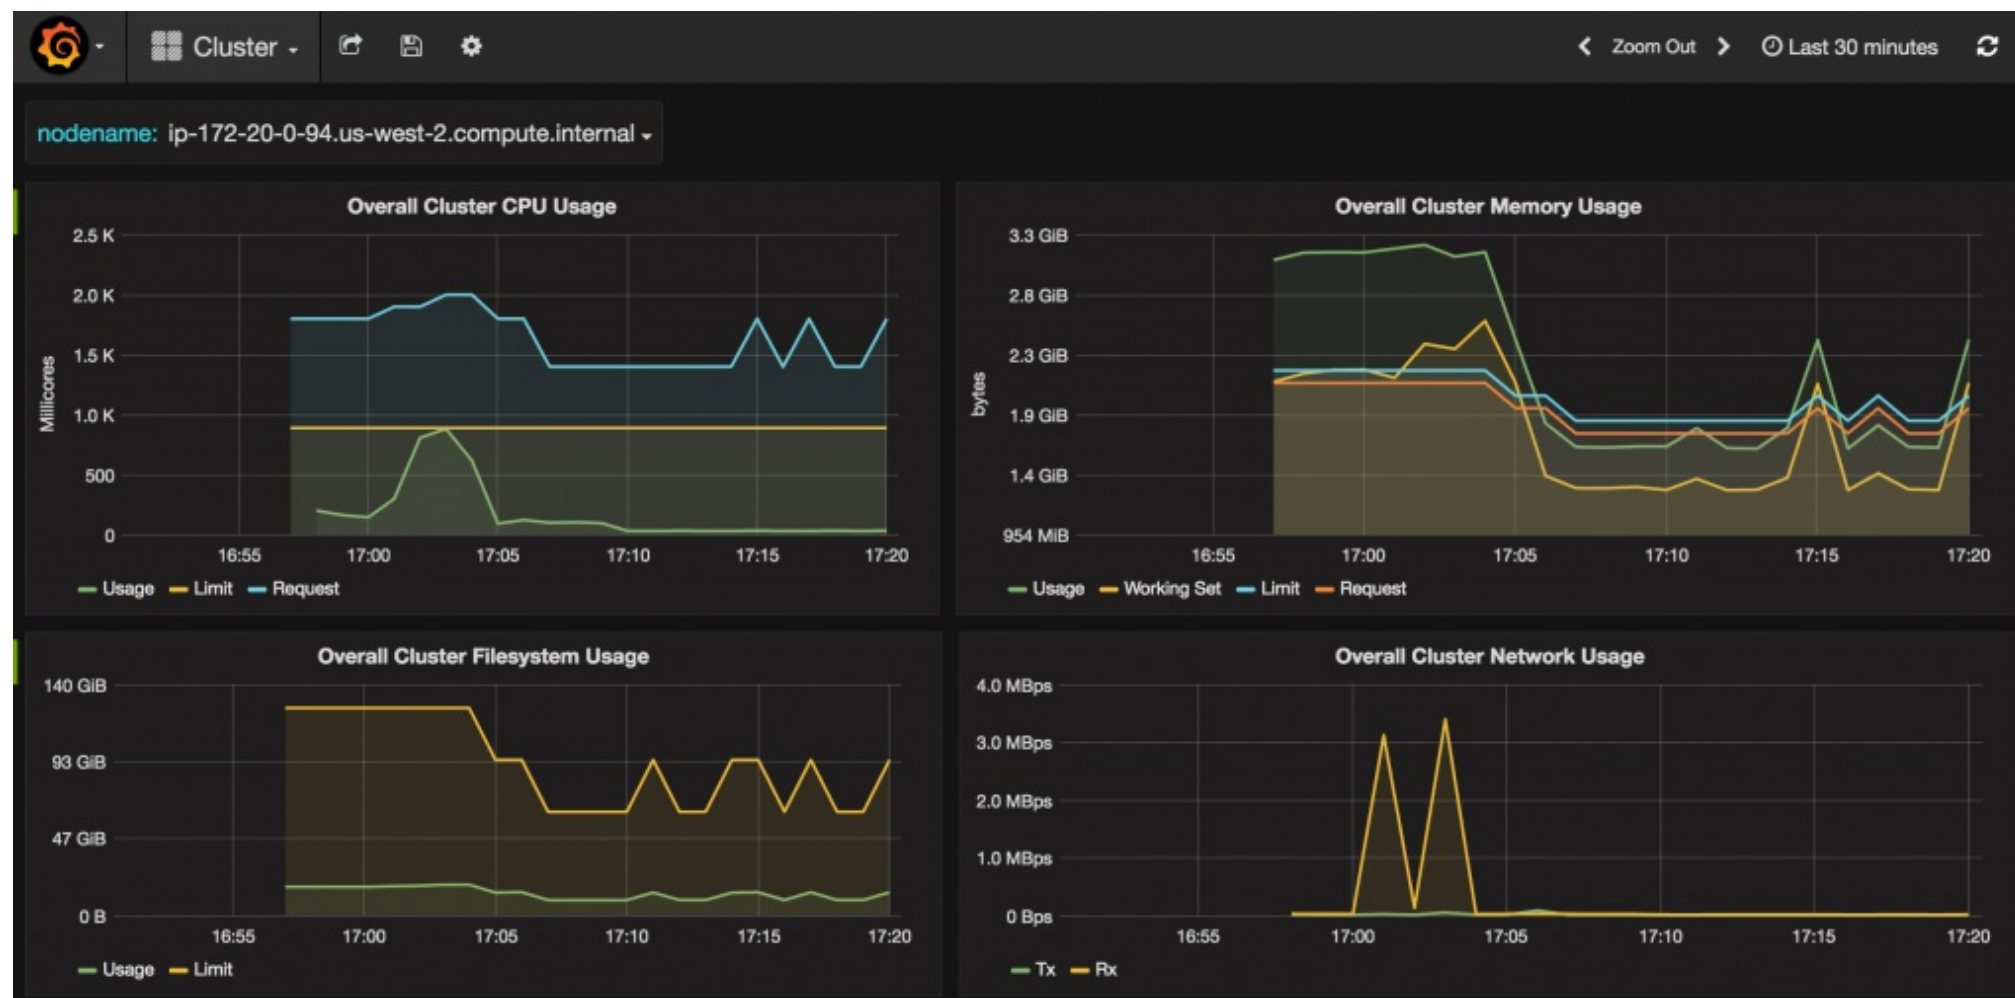
MONITORING

- Liczby
- Trendy
- Zależności

METRYKA


Nazwa	Etykiety	Wartość
traefik_requests_total	code="200", method="GET"	3001

MONITORING



Example from [couchbase blog](#)

MONITORING



JAK ZNALEŹĆ WŁAŚCIWE METRYKI?

JAK ZNALEŹĆ WŁAŚCIWE METRYKI?

- **USE**
- **RED**

USE

Utilization	the average time that the resource was busy servicing work
--------------------	--

Saturation	extra work which it can't service, often queued
-------------------	---

Errors	the count of error events
---------------	---------------------------

Documented and Promoted by [Berdan Gregg](#)

USE

- **Utilization:** as a percent over a time interval: "one disk is running at 90% utilization".
- **Saturation:**
- **Errors:**

USE

- **Utilization:**
- **Saturation:** as a queue length. eg, "the CPUs have an average run queue length of four".
- **Errors:**

USE

- **utilization:**
- **saturation:**
- **errors:** scalar counts. eg, "this network interface drops packages".

USE

- **traditionally** more instance oriented
- still useful in the microservices world

RED

Rate

How busy is your service?

Error

Errors

Duration

What is the latency of my service?

[Tom Wilkie's guideline for instrumenting applications.](#)

RED

- **Rate** - how many request per seconds handled
- **Error**
- **Duration** (distribution)

RED

- **Rate**
- **Error** - how many request per seconds handled we failed
- **Duration**

RED

- **Rate**
- **Error**
- **Duration** - how long the requests took

RED

- Follow Four Golden Signals by Google SREs [1]
- Focus on what matters for end-users

[1] Latency, Traffic, Errors, Saturation ([src](#))

RED

- not recommended for batch-oriented or streaming services

IMPLEMENTACJA Z PROMETHEUS



PROMETHEUS STACK

- Prometheus
- Alertmanager
- Grafana

PROMETHEUS

- Wide support for languages
- Metrics collected over HTTP *metrics/*
- Pull model (see *scrape time*), possible push

METRICS IN TEXT

```
# HELP order_mgmt_audit_duration_seconds Multiprocess metric
# TYPE order_mgmt_audit_duration_seconds summary
order_mgmt_audit_duration_seconds_count{status_code="200"} 41.0
order_mgmt_audit_duration_seconds_sum{status_code="200"} 27.4457
order_mgmt_audit_duration_seconds_count{status_code="500"} 1.0
order_mgmt_audit_duration_seconds_sum{status_code="500"} 0.71663
# HELP order_mgmt_duration_seconds Multiprocess metric
# TYPE order_mgmt_duration_seconds summary
order_mgmt_duration_seconds_count{method="GET",path="/complex",s
order_mgmt_duration_seconds_sum{method="GET",path="/complex",st
order_mgmt_duration_seconds_count{method="GET",path="/",status_co
order_mgmt_duration_seconds_sum{method="GET",path="/",status_coc
order_mgmt_duration_seconds_count{method="GET",path="/complex",s
order_mgmt_duration_seconds_sum{method="GET",path="/complex",st
```


METRICS IN TEXT

```
# HELP go_gc_duration_seconds A summary of the GC invocation duration
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 9.01e-05
go_gc_duration_seconds{quantile="0.25"} 0.000141101
go_gc_duration_seconds{quantile="0.5"} 0.000178902
go_gc_duration_seconds{quantile="0.75"} 0.000226903
go_gc_duration_seconds{quantile="1"} 0.006099658
go_gc_duration_seconds_sum 18.749046756
go_gc_duration_seconds_count 89273
```

PROMETHEUS EXPORTERS

Exporters:

- MongoDB
- Mysql
- Postgresql
- ...
- also Blackbox exporter

[Example](#)

PROMETHEUS PromQL

Powerful query language:

```
histogram_quantile(0.9, rate(http_request_duration_seconds_bucket[10m])  
predict_linear  
rate(http_requests_total{job="api-server"}[5m])  
irate(http_requests_total{job="api-server"}[5m])  
holt_winters()
```

PROMETHEUS PromQL

You can also use it for alarming:

```
ALERT ProductionAppServiceInstanceDown
  IF up { environment = "production", app =~ ".+" } == 0
  FOR 4m
  ANNOTATIONS {
    summary = "Instance of {{$labels.app}} is down",
    description = " Instance  {{$labels.instance}} of app {{$labels.app}}."
  }
```

METRICS

- Counter - just up
- Gauge - up/down
- Summary
- Histogram

HISTOGRAM

traefik_duration_seconds_bucket
{method="GET,code="200"}

{le="0.1"}	2229
{le="0.3"}	107
{le="1.2"}	100
{le="5"}	4
{le="+Inf"}	2
_sum	
_count	2342

SUMMARY

http_request_duration_seconds

{quantile="0.5"}	4
{quantile="0.9"}	5
http_request_duration_seconds_sum	9
http_request_duration_seconds_count	3

HISTOGRAM / SUMMARY:

- Latency of services
- Request or Request size

RED

Metric + PromQL:

```
sum(irate(order_mgmt_duration_seconds_count  
{job=~".*"}[1m])) by (status_code)
```

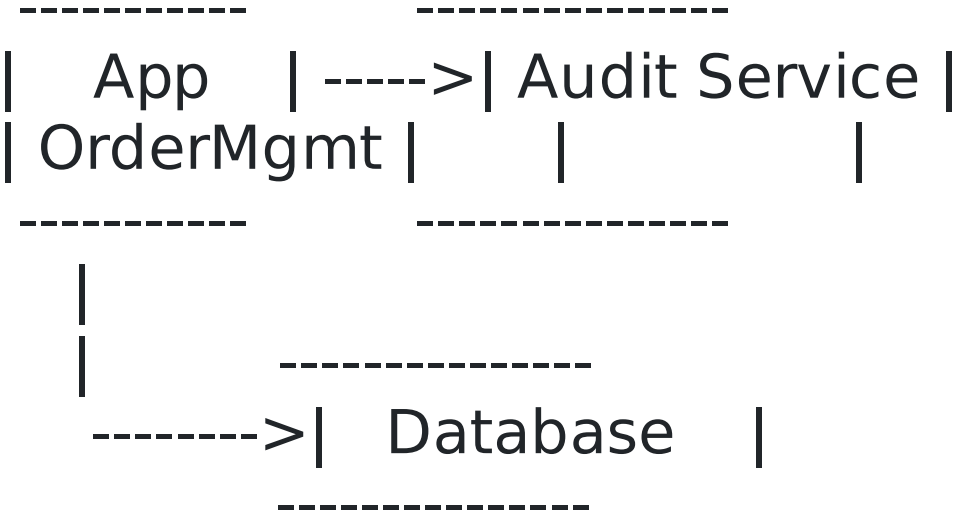
PROMETHEUS + PYTHON



PYTHON CLIENT

- `client_python`
- Counter
- Gauge
- Summary
- Histogram

DEMO: SIMPLE REST SERVICE



DEMO:

- <http://127.0.0.1:8080> - service
- <http://127.0.0.1:8080/metrics/>
- <http://127.0.0.1:9090> - prometheus
- <http://127.0.0.1:3000> - grafana
- <http://127.0.0.1:9093> - alertmanager

DEMO: PYTHON CODE

- Metric Definition
- Metric Collection

DEMO: SIMULATING CALLS

```
curl 127.0.0.1:8080/hello  
curl 127.0.0.1:8080/world  
curl 127.0.0.1:8080/complex
```

DEMO: SIMULATING CALLS

```
curl 127.0.0.1:8080/complex?is_srv_error=True  
curl 127.0.0.1:8080/complex?is_db_error=True  
curl 127.0.0.1:8080/complex?db_sleep=3&srv_sleep=2  
# load generator  
make srv_wrk_random
```


DEMO: PROM STACK

- Prometheus dashboard and config
- AlertManager dashboard and config
- Simulate the successful and failed calls
- Simple Queries for rate


PromQL

```
sum(irate(order_mgmt_duration_seconds_count{job=~".*"}[1m]))  
  by (status_code)
```

PromQL

```
order_mgmt_duration_seconds_sum{job=~".*"} or  
order_mgmt_database_duration_seconds_sum{job=~".*"} or  
order_mgmt_audit_duration_seconds_sum{job=~".*"}
```

MONITORING INGRESS



- --web.metrics.prometheus

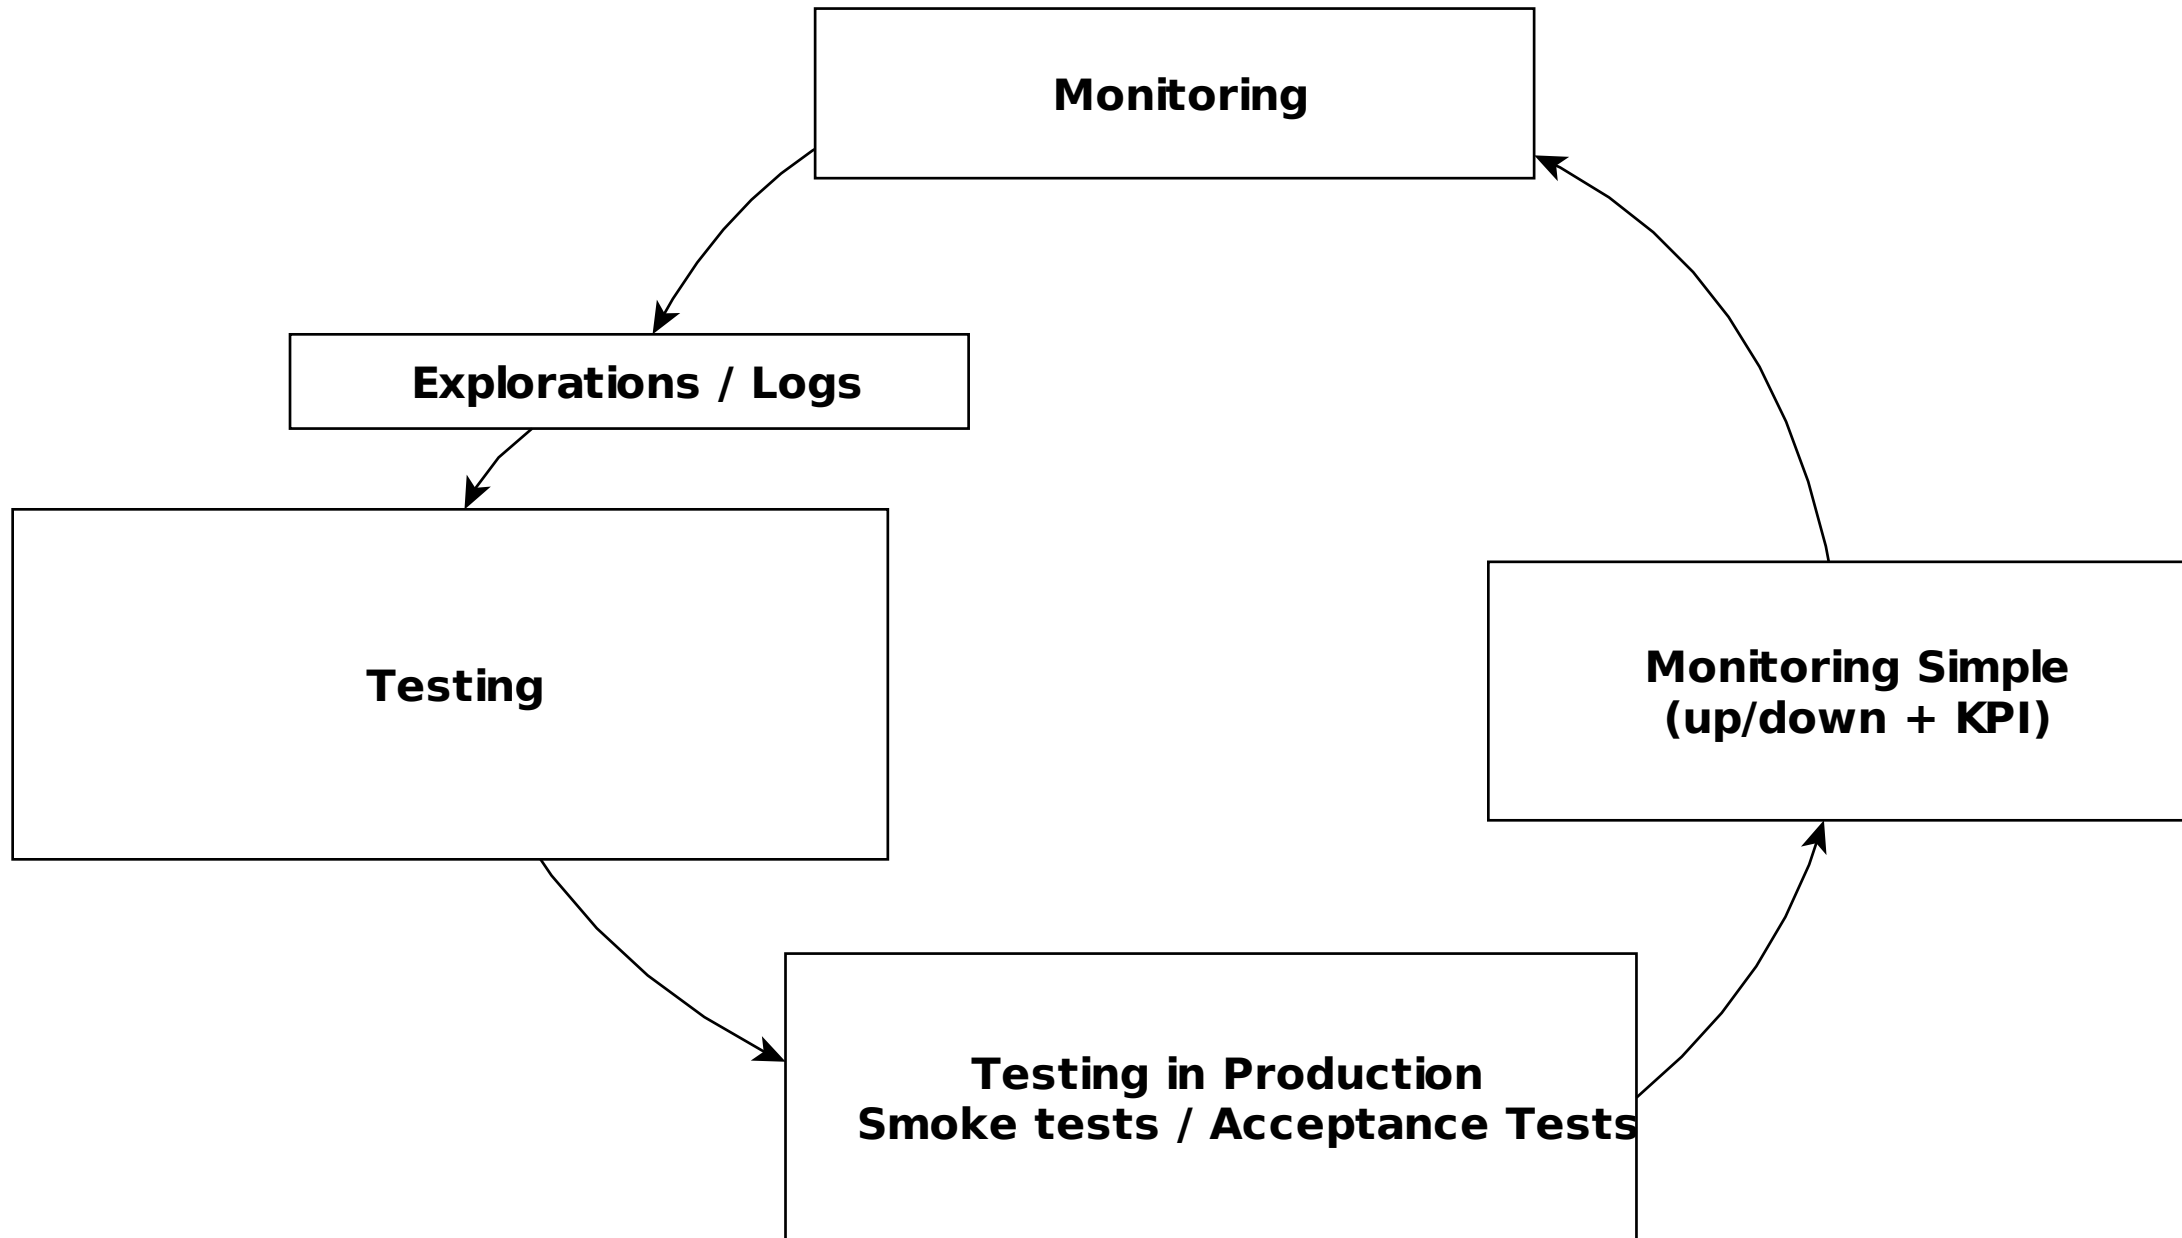
BEST PRACTISES

- Prefix for the metric names is your service name
- Under higher load, you need to have multiprocessing, otherwise your service will hang
- You can start simple (whether sth is up and down), later you can add more complex rules

SUMMARY

- Monitoring saves your time
- Checking logs **Kibana** vs **Grafana** is like debugging vs having tests
- Logging -> high TCO

SUMMARY



THANK YOU

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE
BE WITH YOU.



QUESTIONS?

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE
BE WITH YOU.



SMACCC



Go



PYTORCH



Warsaw Office in BL Astoria:



BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE
BE WITH YOU.



PROMETHUS - LABELS IN ALERT RULES

The labels are propagated to alert rules:

```
ALERT ProductionAppServiceInstanceDown
  IF up { environment = "production", app =~ ".+" } == 0
  FOR 4m
  ANNOTATIONS {
    summary = "Instance of {{$labels.app}} is down",
    description = " Instance  {{$labels.instance}} of app {{$labels.app}} ."
  }
```

see ../src/prometheus/etc/alert.rules

ALERTMANGER - LABELS IN ALERTMANGER

Call somebody if the label is **severity=page**:

```
---
group_by: [cluster]
# If an alert isn't caught by a route, send it to the pager.
receiver: team-pager
routes:
- match:
    severity: page
  receiver: team-pager

receivers:
- name: team-pager
  opsgenie_configs:
  - api_key: $API_KEY
    teams: example_team
```

see ../src/alertmanager/*.conf

PROMETHEUS - PUSH MODEL

- See:

<https://prometheus.io/docs/instrumenting/pushing/>

Good for short living jobs in your cluster.

DESIGNING METRIC NAMES

Which one is better?

- `request_duration{app=my_app}`
- `my_app_request_duration`

DESIGNING METRIC NAMES

Which one is better?

- `order_mgmt_db_duration_seconds_sum`
- `order_mgmt_duration_seconds_sum{dep_name='db'}`

PROMETHEUS + K8S = <3

**LABELS ARE PROPAGATED FROM K8S TO
PROMETHEUS**

INTEGRATION WITH PROMETHEUS

cat memcached-0-service.yaml

```
---
apiVersion: v1
kind: Service
metadata:
  name: memcached-0
  labels:
    app: memcached
    kubernetes.io/name: "memcached"
    role: shard-0
  annotations:
    prometheus.io/scrape: "true"
    prometheus.io/scheme: "http"
    prometheus.io/path: "metrics"
    prometheus.io/port: "9150"
spec:
```

<https://github.com/skarab7/kubernetes-memcached>